

Apache Spark Internals

Pietro Michiardi

Eurecom

Acknowledgments & Sources

● Sources

- ▶ Research papers:
 - ★ <https://spark.apache.org/research.html>
- ▶ Blogs:
 - ★ <https://trongkhoanguyenblog.wordpress.com/>
 - ★ <http://hxquangnhat.com/>

● Contributors

- ▶ Quang-Nhat Hoang-Xuan
- ▶ Khoa Nguyen Trong

Introduction and Motivations

What is Apache Spark

• Project goals

- ▶ Generality: diverse workloads, operators, job sizes
- ▶ Low latency: sub-second
- ▶ Fault tolerance: faults are the norm, not the exception
- ▶ Simplicity: often comes from generality

Motivations

- **Software engineering point of view**

- ▶ Hadoop code base is huge
- ▶ Contributions/Extensions to Hadoop are cumbersome
- ▶ Java-only hinders wide adoption, but Java support is fundamental

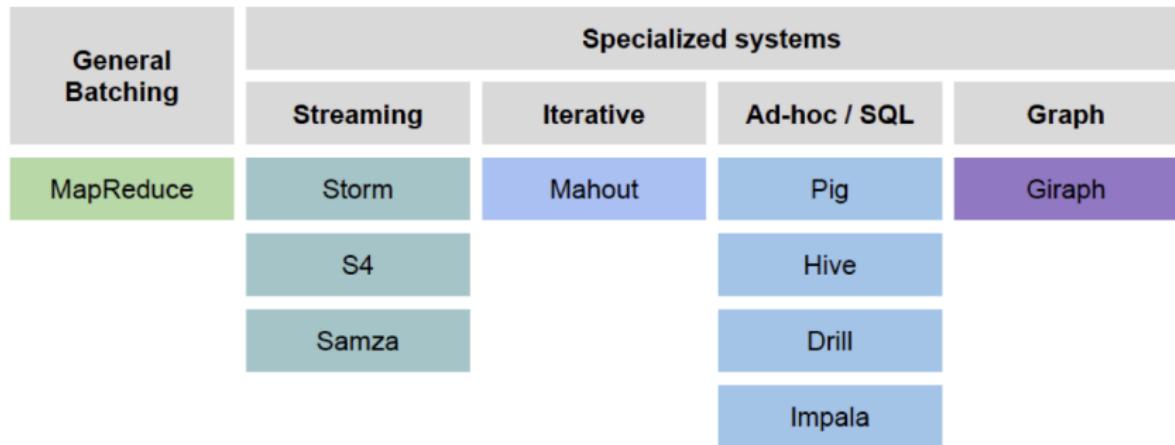
- **System/Framework point of view**

- ▶ Unified pipeline
- ▶ Simplified data flow
- ▶ Faster processing speed

- **Data abstraction point of view**

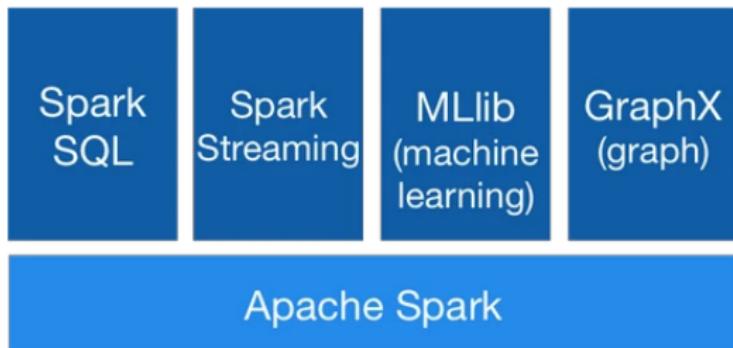
- ▶ New fundamental abstraction RDD
- ▶ Easy to extend with new operators
- ▶ More descriptive computing model

Hadoop: No Unified Vision



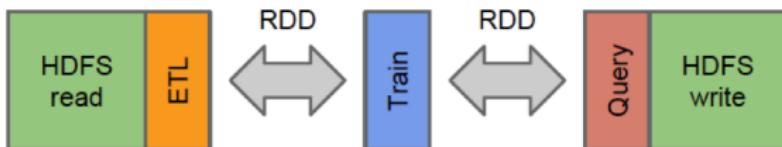
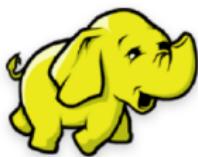
- Sparse modules
- Diversity of APIs
- Higher operational costs

SPARK: A Unified Pipeline



- Spark Streaming (stream processing)
- GraphX (graph processing)
- MLlib (machine learning library)
- Spark SQL (SQL on Spark)

A Simplified Data Flow



Hadoop: Bloated Computing Model

```
public class WordCount {  
  
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
  
    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
  
        public void reduce(Text key, Iterable<IntWritable> values, Context context)  
            throws IOException, InterruptedException {  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            context.write(key, new IntWritable(sum));  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
  
        Job job = new Job(conf, "wordcount");  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.waitForCompletion(true);  
    }  
}
```

SPARK: Descriptive Computing Model

```
val file = sc.textFile("hdfs://...")  
  
val counts = file.flatMap(line => line.split(" ")).  
    .map(word => (word,1))  
    .reduceByKey(_ + _)  
  
counts.saveAsTextFile("hdfs://...")
```

- Organize computation into multiple stages in a processing pipeline
 - ▶ **Transformations** apply user code to distributed data in parallel
 - ▶ **Actions** assemble final output of an algorithm, from distributed data

Faster Processing Speed

- **Let's focus on iterative algorithms**

- ▶ Spark is faster thanks to the simplified data flow
 - ▶ We avoid materializing data on HDFS after each iteration

- **Example: k-means algorithm, 1 iteration**

- ▶ HDFS Read
 - ▶ **Map**(*Assign sample to closest centroid*)
 - ▶ **GroupBy**(Centroid_ID)
 - ▶ NETWORK Shuffle
 - ▶ **Reduce**(*Compute new centroids*)
 - ▶ HDFS Write

Code Base (2012)

Spark core: 16,000 LOC

Operators: 2000

Scheduler: 2500

Block manager: 2700

Networking: 1200

Accumulators: 200

Broadcast: 3500

Interpreter:
3300 LOC

Hadoop I/O:
400 LOC

Mesos backend:
700 LOC

Standalone backend:
1700 LOC

- 2012 (version 0.6.x): 20,000 lines of code
- 2014 (branch-1.0): 50,000 lines of code

Anatomy of a Spark Application

A Very Simple Application Example

```
val sc = new SparkContext("spark://...", "MyJob", home,
jars)

val file = sc.textFile("hdfs://...") // This is an RDD

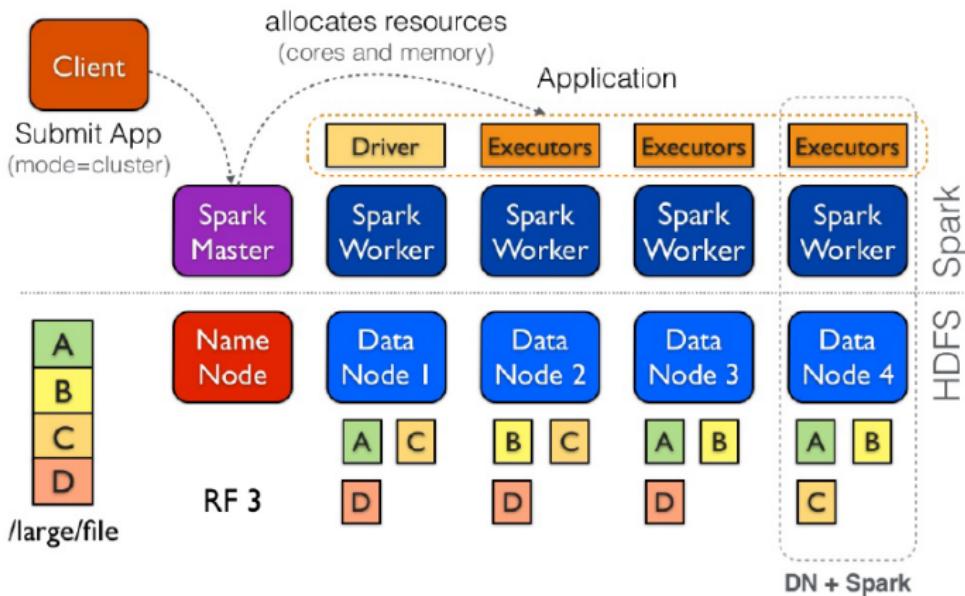
val errors = file.filter(_.contains("ERROR")) // This is
an RDD

errors.cache()

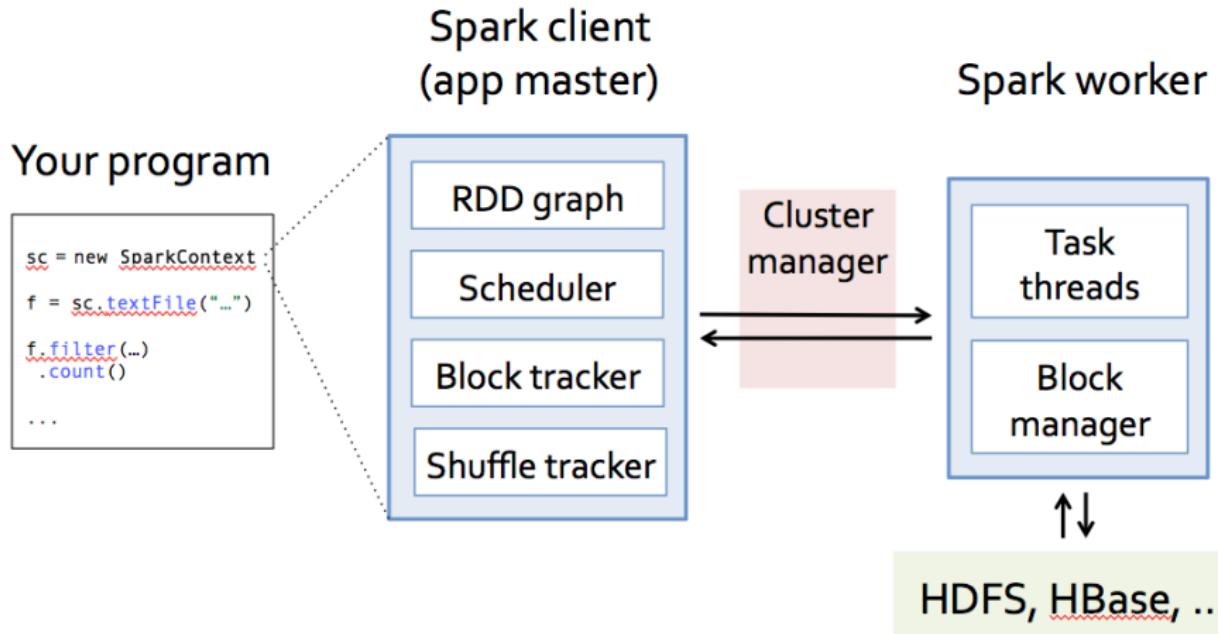
errors.count() // This is an action
```

Spark Applications: The Big Picture

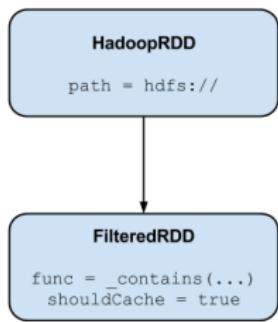
- There are two ways to manipulate data in Spark
 - Use the interactive shell, *i.e.*, the REPL
 - Write standalone applications, *i.e.*, driver programs



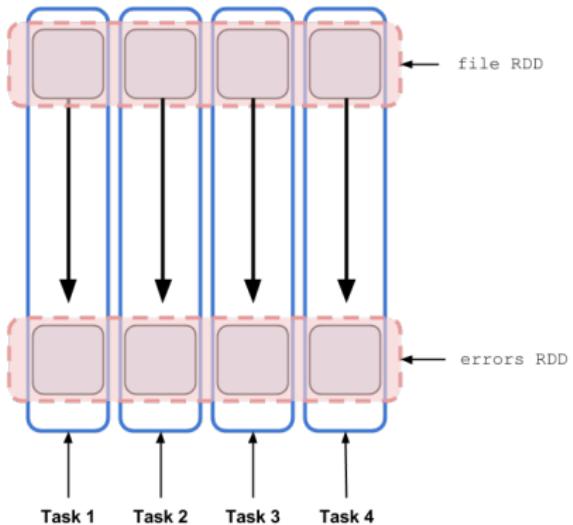
Spark Components: details



The RDD graph: dataset vs. partition views



Worker 1 Worker 2 Worker 3 Worker 4



Data Locality

- **Data locality principle**

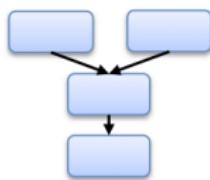
- ▶ Same as for Hadoop MapReduce
- ▶ Avoid network I/O, workers should manage local data

- **Data locality and caching**

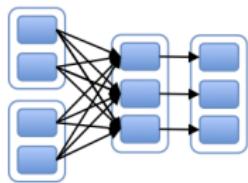
- ▶ First run: data not in cache, so use HadoopRDD's locality prefs (from HDFS)
- ▶ Second run: FilteredRDD is in cache, so use its locations
- ▶ If something falls out of cache, go back to HDFS

Lifetime of a Job in Spark

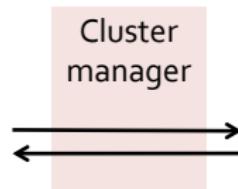
RDD Objects



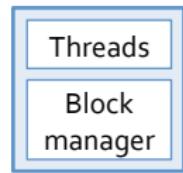
DAG Scheduler



Task Scheduler



Worker



`rdd1.join(rdd2)
.groupBy(...)
.filter(...)`

Build the operator DAG

Split the DAG into
stages of tasks

Submit each stage and
its tasks as ready

Launch tasks via Master

Retry failed and strag-
gler tasks

Execute tasks

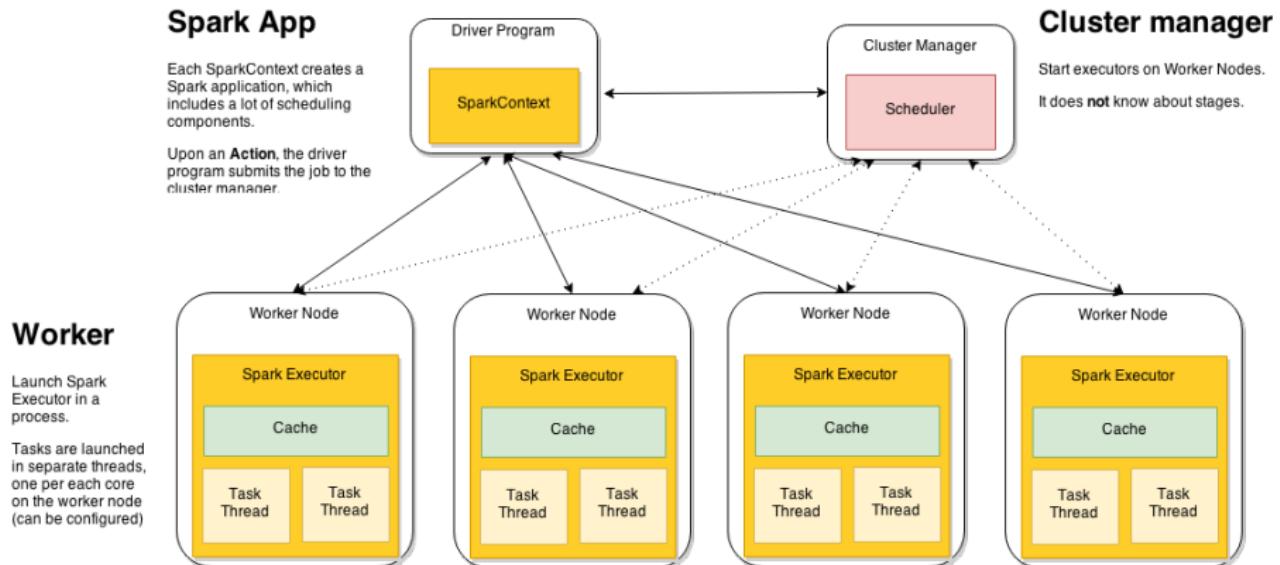
Store and serve blocks

In Summary

- Our example Application: a **jar** file
 - ▶ Creates a `SparkContext`, which is the core component of the driver
 - ▶ Creates an input `RDD`, from a file in HDFS
 - ▶ Manipulates the input `RDD` by applying a `filter(f: T => Boolean)` transformation
 - ▶ Invokes the action `count()` on the transformed `RDD`
- The DAG Scheduler
 - ▶ Gets: `RDDs`, functions to run on each partition and a listener for results
 - ▶ Builds *Stages* of *Tasks* objects (code + preferred location)
 - ▶ Submits Tasks to the **Task Scheduler** as ready
 - ▶ Resubmits failed *Stages*
- The Task Scheduler
 - ▶ Launches *Tasks* on executors
 - ▶ Relaunches failed *Tasks*
 - ▶ Reports to the DAG Scheduler

Spark Deployments

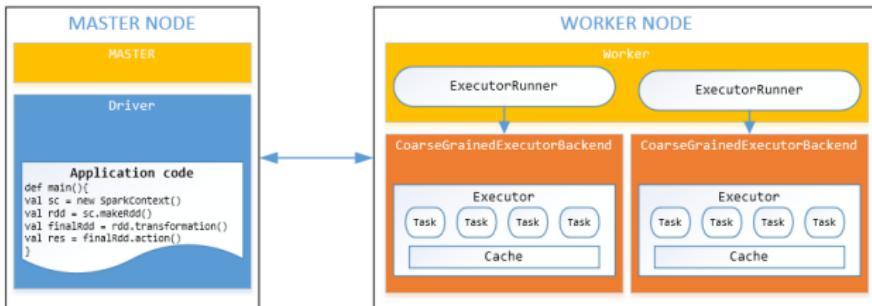
Spark Components: System-level View



Spark Deployment Modes

- **The Spark Framework can adopt several cluster managers**
 - ▶ *Local Mode*
 - ▶ *Standalone mode*
 - ▶ *Apache Mesos*
 - ▶ *Hadoop YARN*
- **General “workflow”**
 - ▶ Spark application creates `SparkContext`, which initializes the `DriverProgram`
 - ▶ Registers to the `ClusterManager`
 - ▶ Ask resources to allocate `Executors`
 - ▶ Schedule Task execution

Worker Nodes and Executors



- **Worker nodes are machines that run executors**
 - ▶ Host one or multiple Workers
 - ▶ One JVM (= 1 UNIX process) per Worker
 - ▶ Each Worker can spawn one or more Executors
- **Executors run tasks**
 - ▶ Run in child JVM (= 1 UNIX process)
 - ▶ Execute one or more task using threads in a ThreadPool

Comparison to Hadoop MapReduce

Hadoop MapReduce

- One Task per UNIX process (JVM), more if JVM reuse
 - MultiThreadedMapper, advanced feature to have threads in Map Tasks
- **Short-lived Executor**, with one **large Task**

Spark

- Tasks run in one or more Threads, within a single UNIX process (JVM)
 - Executor process statically allocated to worker, even with no threads
- **Long-lived Executor**, with many **small Tasks**

Benefits of the Spark Architecture

- **Isolation**

- ▶ Applications are completely isolated
- ▶ Task scheduling *per application*

- **Low-overhead**

- ▶ Task setup cost is that of spawning a thread, not a process
- ▶ 10-100 times faster
- ▶ **Small tasks → mitigate effects of data skew**

- **Sharing data**

- ▶ Applications cannot share data in memory natively
- ▶ Use an external storage service like Tachyon

- **Resource allocation**

- ▶ Static process provisioning for executors, even without active tasks
- ▶ Dynamic provisioning under development

Resilient Distributed Datasets

M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauley, M.J. Franklin, S. Shenker, I. Stoica.

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing,

USENIX Symposium on Networked Systems Design and Implementation, 2012

What is an RDD

- **RDD are partitioned, locality aware, distributed collections**
 - ▶ RDD are *immutable*
- **RDD are data structures that:**
 - ▶ Either point to a direct data source (e.g. HDFS)
 - ▶ Apply some transformations to its parent RDD(s) to generate new data elements
- **Computations on RDDs**
 - ▶ Represented by *lazily evaluated* lineage DAGs composed by chained RDDs

RDD Abstraction

- **Overall objective**

- ▶ Support a wide array of operators (more than just Map and Reduce)
- ▶ Allow arbitrary composition of such operators

- **Simplify scheduling**

- ▶ Avoid to modify the scheduler for each operator

→ The question is: *How to capture dependencies in a general way?*

RDD Interfaces

- **Set of partitions (“splits”)**
 - ▶ Much like in Hadoop MapReduce, each RDD is associated to (input) partitions
- **List of dependencies on parent RDDs**
 - ▶ This is completely new w.r.t. Hadoop MapReduce
- **Function to compute a partition given parents**
 - ▶ This is actually the “user-defined code” we referred to when discussing about the Mapper and Reducer classes in Hadoop
- **Optional preferred locations**
 - ▶ This is to enforce data locality
- **Optional partitioning info (Partitioner)**
 - ▶ This really helps in some “advanced” scenarios in which you want to pay attention to the behavior of the shuffle mechanism

Hadoop RDD

- `partitions` = one per HDFS block
- `dependencies` = none
- `compute(partition)` = read corresponding block
- `preferredLocations(part)` = HDFS block location
- `partitioner` = none

Filtered RDD

- `partitions` = same as parent RDD
- `dependencies` = *one-to-one* on parent
- `compute(partition)` = compute parent and filter it
- `preferredLocations(part)` = none (*ask parent*)
- `partitioner` = none

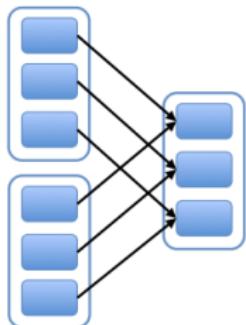
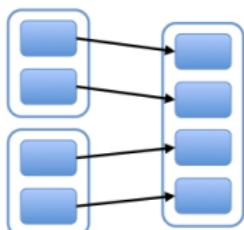
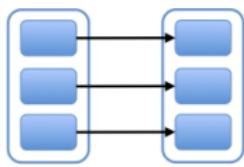
Joined RDD

- **partitions** = one per reduce task
- **dependencies** = *shuffle* on *each* parent
- **compute(partition)** = read and join shuffled data
- **preferredLocations(part)** = none
- **partitioner** = HashPartitioner(numTask)¹

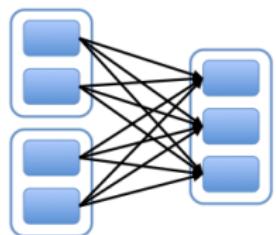
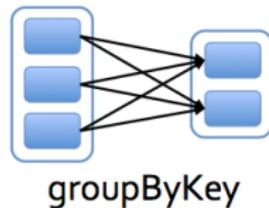
¹Spark knows this data is hashed.

Dependency Types (1)

Narrow dependencies



Wide dependencies



Dependency Types (2)

- **Narrow dependencies**

- ▶ Each partition of the parent RDD is used by at most one partition of the child RDD
- ▶ Task can be executed locally and we don't have to shuffle. (Eg: map, flatMap, filter, sample)

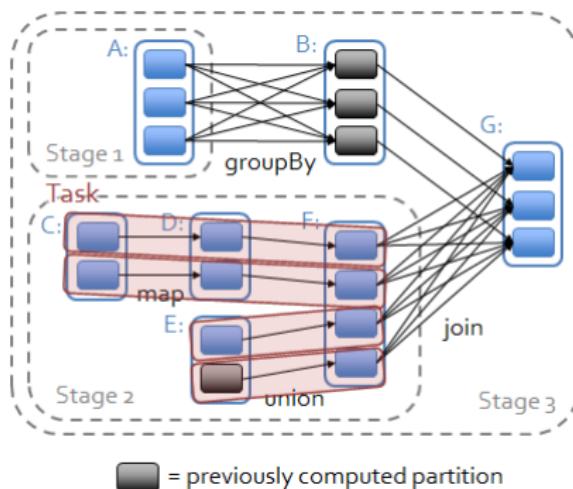
- **Wide Dependencies**

- ▶ Multiple child partitions may depend on one partition of the parent RDD
- ▶ This means we have to shuffle data **unless the parents are hash-partitioned** (Eg: sortByKey, reduceByKey, groupByKey, cogroupByKey, join, cartesian)

Dependency Types: Optimizations

- Benefits of Lazy evaluation

- ▶ The DAG Scheduler optimizes *Stages* and *Tasks* before submitting them to the Task Scheduler
- ▶ **Piplining** narrow dependencies within a Stage
- ▶ **Join plan selection** based on partitioning
- ▶ **Cache reuse**



Operations on RDDs: Transformations

• Transformations

- ▶ Set of operations on a RDD that define how they should be transformed
- ▶ As in relational algebra, the application of a transformation to an RDD yields a new RDD (because RDD are *immutable*)
- ▶ Transformations are lazily evaluated, which allow for optimizations to take place before execution

• Examples (not exhaustive)

- ▶ `map(func)`, `flatMap(func)`, `filter(func)`
- ▶ `groupByKey()`
- ▶ `reduceByKey(func)`, `mapValues(func)`, `distinct()`,
`sortByKey(func)`
- ▶ `join(other)`, `union(other)`
- ▶ `sample()`

Operations on RDDs: Actions

- **Actions**

- ▶ Apply transformation chains on RDDs, eventually performing some additional operations (e.g., counting)
- ▶ Some actions only store data to an external data source (e.g. HDFS), others fetch data from the RDD (and its transformation chain) upon which the action is applied, and convey it to the driver

- **Examples** (not exhaustive)

- ▶ `reduce(func)`
- ▶ `collect()`, `first()`, `take()`, `foreach(func)`
- ▶ `count()`, `countByKey()`
- ▶ `saveAsTextFile()`

Operations on RDDs: Final Notes

- **Look at return types!**

- ▶ Return type: $\text{RDD} \rightarrow \text{transformation}$
- ▶ Return type: built-in scala/java types such as `int`, `long`,
`List<Object>`, `Array<Object>` → `action`

- **Caching is a transformation**

- ▶ Hints to keep RDD in memory after its first evaluation

- **Transformations depend on RDD “flavor”**

- ▶ `PairRDD`
- ▶ `SchemaRDD`

RDD Code Snippet

```
abstract class RDD[T: ClassTag]{
    @transient private var sc: SparkContext,
    @transient private var deps: Seq[Dependency[_]]
} extends Serializable with Logging {
```

- **SparkContext**

- ▶ This is the main entity responsible for setting up a job
- ▶ Contains SparkConfig, Scheduler, entry point of running jobs (runJobs)

- **Dependencies**

- ▶ Input RDD(s)

RDD.map operation Snippet

- **Map: RDD[T] → RDD[U]**

```
/**  
 * Return a new RDD by applying a function to all elements of this RDD.  
 */  
def map[U: ClassTag](f: T => U): RDD[U] = new MappedRDD(this, sc.clean(f))
```

- **MappedRDD**

- ▶ For each element in a partition, apply function f

```
private[spark]  
class MappedRDD[U: ClassTag, T: ClassTag](prev: RDD[T], f: T => U)  
  extends RDD[U](prev) {  
  
  override def getPartitions: Array[Partition] = firstParent[T].partitions  
  
  override def compute(split: Partition, context: TaskContext) =  
    firstParent[T].iterator(split, context).map(f)  
}
```

RDD Iterator Code Snipped

- Method to go through an RDD and apply function f
 - First, check local cache
 - If not found, compute the RDD

```
/**  
 * Internal method to this RDD; will read from cache if applicable, or otherwise compute it.  
 * This should 'not' be called by users directly, but is available for implementors of custom  
 * subclasses of RDD.  
 */  
final def iterator(split: Partition, context: TaskContext): Iterator[T] = {  
    if (storageLevel != StorageLevel.NONE) {  
        SparkEnv.get.cacheManager.getOrCompute(this, split, context, storageLevel)  
    } else {  
        computeOrReadCheckpoint(split, context)  
    }  
}
```

Storage Levels

- Disk
- Memory
- Off Heap (e.g. external memory stores like Tachyon)
- De-serialized

Making RDD from local collections

- Convert a local (on the driver) Seq[T] into RDD[T]

```
// Methods for creating RDDs

/** Distribute a local Scala collection to form an RDD. */
def parallelize[T: ClassTag](seq: Seq[T], numSlices: Int = defaultParallelism): RDD[T] = {
  new ParallelCollectionRDD[T](this, seq, numSlices, Map[Int, Seq[String]]())
}

/** Distribute a local Scala collection to form an RDD. */
def makeRDD[T: ClassTag](seq: Seq[T], numSlices: Int = defaultParallelism): RDD[T] = {
  parallelize(seq, numSlices)
}

/** Distribute a local Scala collection to form an RDD, with one or more
 * location preferences (hostnames of Spark nodes) for each object.
 * Create a new partition for each collection item. */
def makeRDD[T: ClassTag](seq: Seq[(T, Seq[String])]): RDD[T] = {
  val indexToPrefs = seq.zipWithIndex.map(t => (t._2, t._1._2)).toMap
  new ParallelCollectionRDD[T](this, seq.map(_._1), seq.size, indexToPrefs)
}
```

Hadoop RDD Code Snippet

- Reading HDFS data as <key, value> records

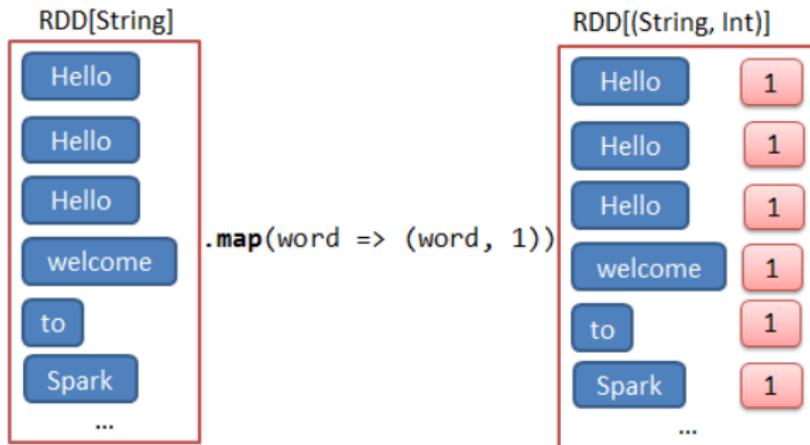
```
/** Get an RDD for a Hadoop file with an arbitrary InputFormat
 *
 * ***Note:*** Because Hadoop's RecordReader class re-uses the same Writable object for each
 * record, directly caching the returned RDD will create many references to the same object.
 * If you plan to directly cache Hadoop writable objects, you should first copy them using
 * a `map` function.
 */
def hadoopFile[K, V](
    path: String,
    inputFormatClass: Class[_ <: InputFormat[K, V]],
    keyClass: Class[K],
    valueClass: Class[V],
    minPartitions: Int = defaultMinPartitions
): RDD[(K, V)] = {
    // A Hadoop configuration can be about 10 KB, which is pretty big, so broadcast it.
    val confBroadcast = broadcast(new SerializableWritable(hadoopConfiguration))
    val setInputPathsFunc = (jobConf: JobConf) => FileInputFormat.setInputPaths(jobConf, path)
    new HadoopRDD(
        this,
        confBroadcast,
        Some(setInputPathsFunc),
        inputFormatClass,
        keyClass,
        valueClass,
        minPartitions
    )
}
```

Understanding RDD Operations

Common Transformations

map (f: T => U)

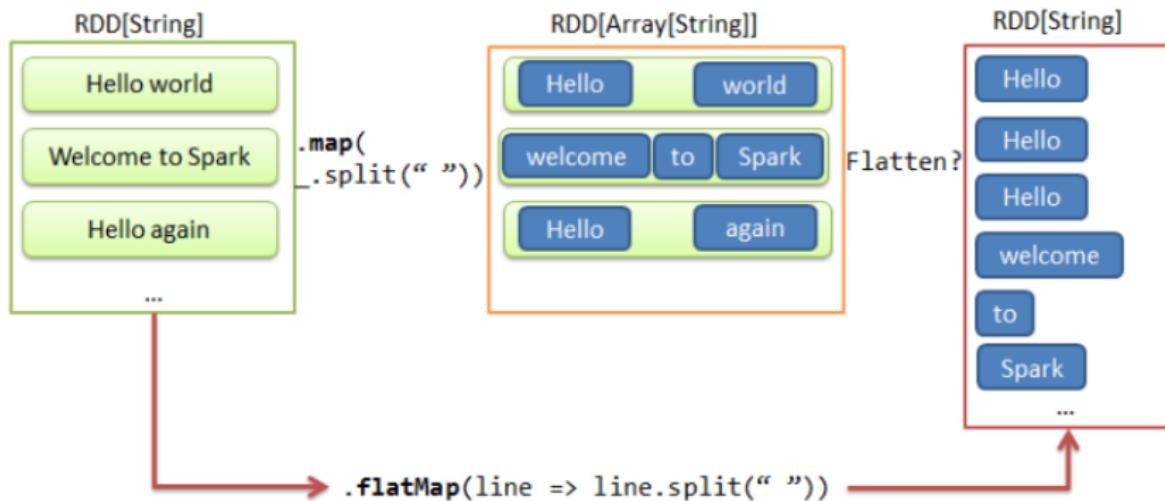
Returns a MappedRDD [U] by applying *f* to each element



Common Transformations

```
flatMap(f: T =>
TraversableOnce[U])
```

Returns a
FlatMappedRDD [U] by
first applying f to each
element, then flattening the
results



Detailed Example: Word Count

Spark Word Count: the driver

```
import org.apache.spark.SparkContext  
  
import org.apache.spark.SparkContext._  
  
val sc = new SparkContext("spark://...", "MyJob", "spark  
home", "additional jars")
```

• Driver and **SparkContext**

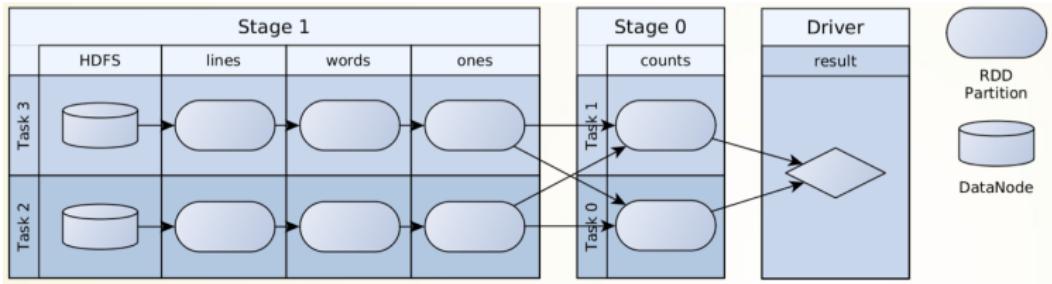
- ▶ A `SparkContext` initializes the application driver, the latter then registers the application to the cluster manager, and gets a list of executors
- ▶ Then, the driver takes full control of the Spark job

Spark Word Count: the code

```
1 val lines = sc.textFile("input")
2 val words = lines.flatMap(_.split(" "))
3 val ones = words.map(_ -> 1)
4 val counts = ones.reduceByKey(_ + _)
5 val result = counts.collectAsMap()
```

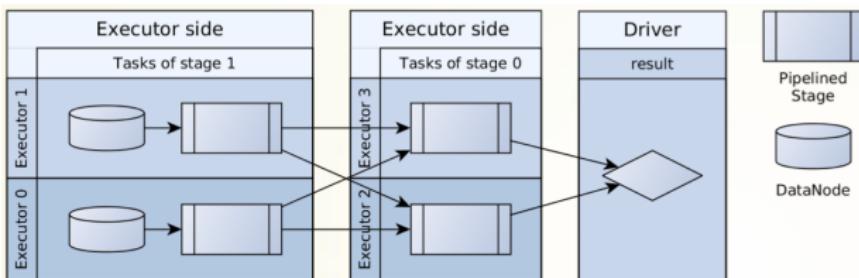
- **RDD lineage DAG is built on driver side with**
 - ▶ Data source RDD(s)
 - ▶ Transformation RDD(s), which are created by transformations
- **Job submission**
 - ▶ An *action* triggers the DAG scheduler to submit a job

Spark Word Count: the DAG



- **Directed Acyclic Graph**
 - ▶ Built from the RDD lineage
- **DAG scheduler**
 - ▶ Transforms the DAG into stages and turns each partition of a stage into a single task
 - ▶ Decides what to run

Spark Word Count: the execution plan



Spark Tasks

- ▶ Serialized RDD lineage DAG + closures of transformations
- ▶ Run by Spark executors

Task scheduling

- ▶ The driver side task scheduler launches tasks on executors according to resource and locality constraints
- ▶ The task scheduler decides where to run tasks

Spark Word Count: the Shuffle phase

```
1 val lines = sc.textFile("input")
2 val words = lines.flatMap(_.split(" "))
3 val ones = words.map(_ -> 1)
4 val counts = ones.reduceByKey(_ + _)
5 val result = counts.collectAsMap()
```

- **reduceByKey transformation**

- ▶ Induces the shuffle phase
- ▶ In particular, we have a *wide dependency*
- ▶ Like in Hadoop MapReduce, intermediate <key,value> pairs are stored on the local file system

- **Automatic combiners!**

- ▶ The `reduceByKey` transformation implements map-side combiners to pre-aggregate data

Caching and Storage

Spark's Storage Module

- **The storage module**

- ▶ Access (I/O) “external” data sources: HDFS, Local Disk, RAM, remote data access through the network
- ▶ Caches RDDs using a variety of “storage levels”

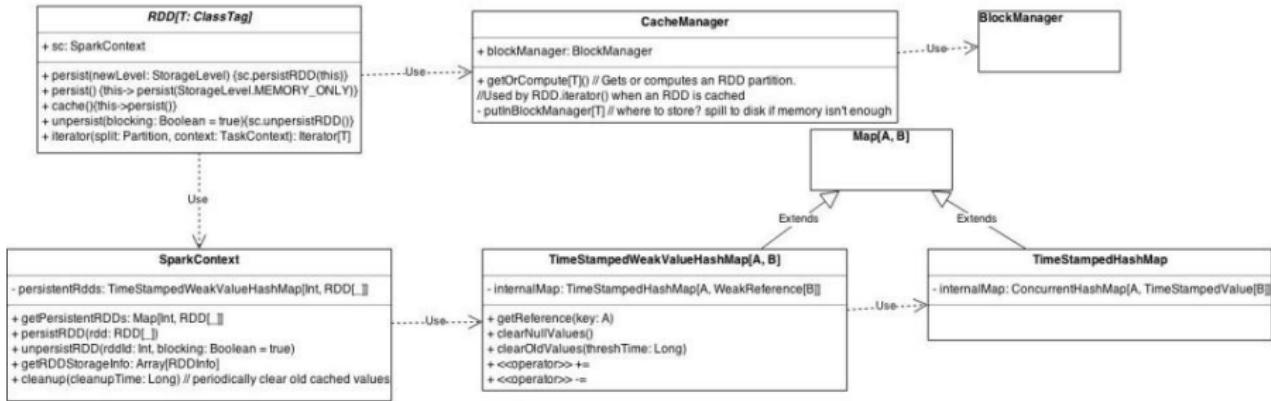
- **Main components**

- ▶ The Cache Manager: uses the Block Manager to perform caching
- ▶ The Block Manager: distributed key/value store

Class Diagram of the Caching Component

RDD Caching flow in Spark-core

Class diagram



How Caching Works

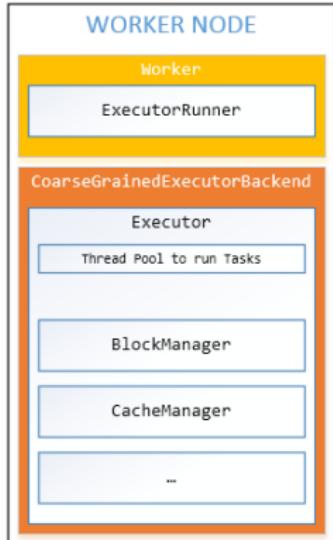
- **Frequently used RDD can be stored in memory**

- ▶ Deciding which RDD to cache is an art!
 - ▶ One method, one short-cut: `persist()`, `cache()`

- **SparkContext keeps track of cached RDD**

- ▶ Uses a data-structred called `persistentRDD`
 - ▶ Maintains references to cached RDD, and eventually call the garbage collector
 - ▶ Time-stamp based invalidation using
`TimeStampedWeakValueHashMap[A, B]`

How Caching Works



```

- rdd.iterator(split:Partition)
-- cacheManager.getOrCompute(rdd, partition)
--- val key = RDDBlockId(rdd.id, partition.index)
--- switch(blockManager.get(key))
---   case Some: return
---   case None:
        val computedValue = rdd.computeOrReadCheckPoint(partition)
        if isRunningLocally
            return computedValue
        else
            cachedValue = putInBlockManager(key, computedValue, storageLevel)
    
```

BlockManager.scala

```

// write-once key-value
// serves both cachedRdds and shuffle data
// tracks storage Level of each block
// can drop data on disk if RAM mem is low
// can replicate data across nodes
def get(blockId){
    tryGetLocal()
    tryGetRemote()
    if no, return None
}
    
```

The Block Manager

- “Write-once” key-value store
 - ▶ One node per worker
 - ▶ No updates, data is immutable

- Main tasks
 - ▶ Serves shuffle data (local or remote connections) and cached RDDs
 - ▶ Tracks the “Storage Level” (RAM, disk) for each block
 - ▶ Spills data to disk if memory is insufficient
 - ▶ Handles data replication, if required

Storage Levels

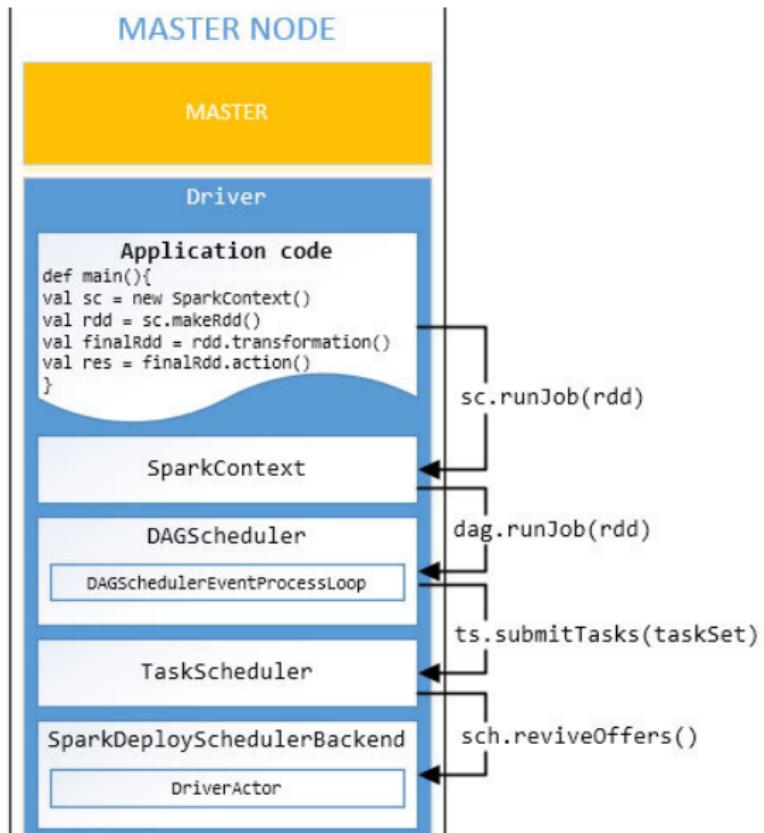
- **The Block Manager can hold data in various storage tiers**
 - ▶ `org.apache.spark.storage.StorageLevel` contains flags to indicate which tier to use
 - ▶ Manual configuration, in the application
 - ▶ Deciding the storage level to use for RDDs is not trivial
- **Available storage tiers**
 - ▶ RAM (default option): if the the RDD doesn't fit in memory, some partitions will not be cached (will be re-computed when needed)
 - ▶ Tachyon (off java heap): reduces garbage collection overhead, the crash of an executor no longer leads to cached data loss
 - ▶ Disk
- **Data format**
 - ▶ Serialized or as Java objects
 - ▶ Replicated partitions

Resource Allocation: Spark Schedulers

Spark Schedulers

- **Two main scheduler components, executed by the driver**
 - ▶ The DAG scheduler
 - ▶ The Task scheduler
- **Objectives**
 - ▶ Gain a broad understanding of how Spark submits Applications
 - ▶ Understand how *Stages* and *Tasks* are built, and their optimization
 - ▶ Understand interaction among various other Spark components

Submitting a Spark Application: A Walk Through



Submitting a Spark Application: Details

```

1 - finalRdd.action()
2 -- sc.runJob()
3 --- dagScheduler.runJob()
4 ---- dagScheduler.submitJob()
5 ----- DAGSchedulerEventProcessLoop.post(JobSubmitted)

```



```

6 - DAGSchedulerEventProcessLoop.onReceive(JobSubmitted)
7 -- dagScheduler.handleJobSubmitted()
8 --- finalStage = new Stage()
9 --- submitStage(finalStage)
10 ---- missingStages = getMissingParentStages(stage).sortById()
11 ---- if missingStages != Nil
12       foreach parent <- missingStages: submitStage(parent)
13 ---- else
14       dagScheduler.submitMissingTasks(stage)

15 def getMissingParentStages(stage){
16   case ShuffleDependency: missing += new Stage()
17   case NarrowDependency: waiting4Visit.push()
18 }


```

```

19 def submitMissingTasks(stage){
20   val loc = getPreferredLocs(stage.rdd)
21   if stage.isShuffleMap
22     tasks:Seq[Task] = partitions.foreach(yield new ShuffleMapTask(p, loc))
23   else
24     tasks:Seq[Task] = partitions.foreach(yield new ResultTask())
25
26 taskScheduler.submitTasks(new TaskSet(tasks))
27 }

```

TaskSchedulerImpl.scala

```

27 def submitTasks(taskSet){
28   schedulableBuilder.addTaskSetManager(new TaskSetManager(taskSet))
29   if (!runInLocal && !hasReceivedTask){
30     starvationTimer.scheduleAtFixedRate(new TimerTask(){...}, STARVATION_DELAY)
31   }
32   schedulerBackend.reviveOffers()
33 }

```

```

34 def sparkDeploySchedulerBackend.reviveOffers(){
35   driverActor ! ReviveOffers
36 }

```

CoarseGrainedSchedulerBackend.scala

```

37 def receiveWithLogging(){
38   case ReviveOffers => makeOffers(){
39     launchTasks(tasks){
40       foreach task <- tasks:
41         executorActor(task.executorId) ! LaunchTask(task)
42     }
43   }
44 }

```

The DAG Scheduler

- **Stage-oriented scheduling**

- Computes a DAG of stages for each job in the application
Lines 10–14, details in Lines 15–27
- Keeps track of which RDD and stage output are materialized
- Determines an optimal schedule, minimizing stages
- Submit stages as sets of Tasks (`TaskSets`) to the Task scheduler
Line 26

- **Data locality principle**

- Uses “preferred location” information (optionally) attached to each RDD
Line 20
- Package this information into Tasks and send it to the Task scheduler

- **Manages Stage failures**

- Failure type: (intermediate) data loss of shuffle output files
- Failed stages will be resubmitted
- NOTE: Task failures are handled by the Task scheduler, which simply resubmit them if they can be computed with no dependency on previous output

The DAG Scheduler: Implementation Details

- **Implemented as an event queue**

- ▶ Uses a daemon thread to handle various kinds of events
Line 6
- ▶ JobSubmitted, JobCancelled, CompletionEvent
- ▶ The thread “swipes” the queue, and routes event to the corresponding handlers

- **What happens when a job is submitted to the DAGScheduler?**

- ▶ JobWaiter object is created
- ▶ JobSubmitted event is fired
- ▶ The daemon thread blocks and wait for a job result
Lines 3, 4

The DAG Scheduler: Implementation Details (2)

- Who handles the JobSubmitted event?

- ▶ Specific handler called `handleJobSubmitted`
Line 6

- Walk-through to the Job Submitted handler

- ▶ Create a new job, called `ActiveJob`
 - ▶ New job starts with only 1 stage, corresponding to the last stage of the job upon which an action is called
Lines 8–9
 - ▶ Use the dependency information to produce additional stages
 - ★ Shuffle Dependency: create a new map stage
Line 16
 - ★ Narrow Dependency: pipes them into a single stage
`getMissingParentStages`

More About Stages

- **What is a DAG**

- ▶ Directed acyclic graph of stages
- ▶ Stage boundaries determined by the shuffle phase
- ▶ Stages are run in *topological order*

- **Definition of a Stage**

- ▶ Set of *independent* tasks
- ▶ All tasks of a stage apply the same function
- ▶ All tasks of a stage have the same dependency type
- ▶ All tasks in a stage belong to a TaskSet

- **Stage types**

- ▶ Shuffle Map Stage: stage tasks results are inputs for another stage
- ▶ Result Stage: tasks compute the final action that initiated a job
(e.g., `count()`, `save()`, etc.)

The Task Scheduler

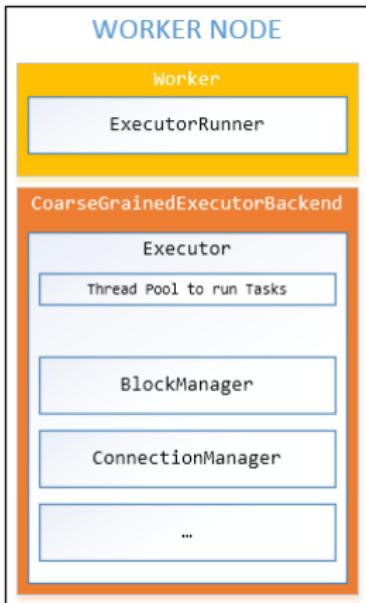
- **Task oriented scheduling**

- ▶ Schedules tasks for a *single* SparkContext
- ▶ Submits tasks sets produced by the DAG Scheduler
- ▶ Retries failed tasks
- ▶ Takes care of *stragglers* with speculative execution
- ▶ Produces events for the DAG Scheduler

- **Implementation details**

- ▶ The Task scheduler creates a `TaskSetManager` to wrap the `TaskSet` from the DAG scheduler
Line 28
- ▶ The `TaskSetManager` class operates as follows:
 - ★ Keeps track of each task status
 - ★ Retries failed tasks
 - ★ Imposes data locality using *delayed scheduling*
Lines 29, 30
- ▶ Message passing implemented using *Actors*, and precisely using the *Akka framework*

Running Tasks on Executors



```

Executor.scala
1 def launchTask(serializedTask){
2     val tr = new TaskRunner(serializedTask)
3     threadPool.execute(tr)
4 }

TaskRunner.scala
5 def run(){
6     executorBackend.statusUpdate(RUNNING)
7     val task = serializer.deserialize(serializedTask)
8     val value = task.run()
9     val res = new DirectTaskResult(serializer.serialize(value))
10    executorBackend.statusUpdate(FINISHED)
11    if(res.size > akkaFrameSize
12        blockManager.putBytes(taskId, res)
13    else
14        return res
15 }

16 task.run()
17 if task is ResultTask
18     val (rdd, func) = ser.deserialize(RDD, taskContext)
19     return result = func(rdd.iterator(partition))
20
21 if task is ShuffleMapTask
22     val (rdd, dep) = ser.deserialize(RDD, ShuffleDependency, taskContext)
23     shuffleWriter = shuffleManager.getWriter(dep.shuffleHandler, partitionId)
24     shuffleWriter.write(rdd.iterator(partition))
25     return shuffleWriter.stop().get()

ShuffleRDD.scala / CoGroupedRDD.scala
26 def compute(split, taskContext){
27     shuffleManager.getReader(dep.shuffleHandler, split.index, taskContext).read()
28 }

```

Running Tasks on Executors

- **Executors run two kinds of tasks**

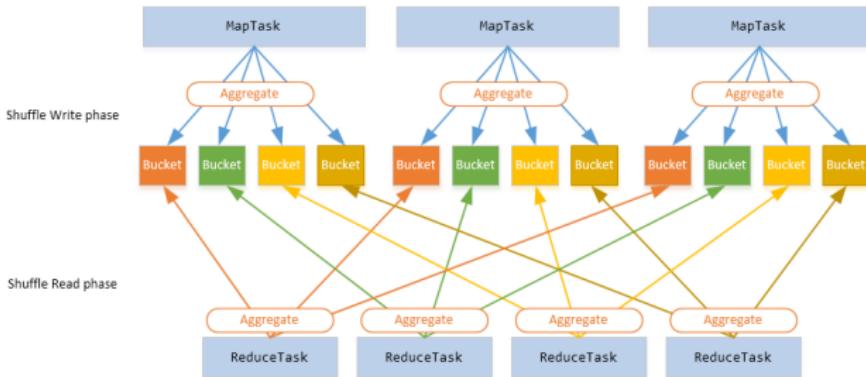
- ▶ ResultTask: apply the action on the RDD, once it has been computed, alongside all its dependencies
Line 19
- ▶ ShuffleTask: use the Block Manager to store shuffle output using the ShuffleWriter
Lines 23, 24
- ▶ The ShuffleRead component depends on the type of the RDD, which is determined by the compute function and the transformation applied to it

Data Shuffling

The Spark Shuffle Mechanism

- Same concept as for Hadoop MapReduce, involving:
 - ▶ Storage of “intermediate” results on the local file-system
 - ▶ Partitioning of “intermediate” data
 - ▶ Serialization / De-serialization
 - ▶ Pulling data over the network
- Transformations requiring a shuffle phase
 - ▶ `groupByKey()`, `reduceByKey()`, `sortByKey()`, `distinct()`
- Various types of Shuffle
 - ▶ *Hash Shuffle*
 - ▶ *Consolidate Hash Shuffle*
 - ▶ *Sort-based Shuffle*

The Spark Shuffle Mechanism: an Illustration

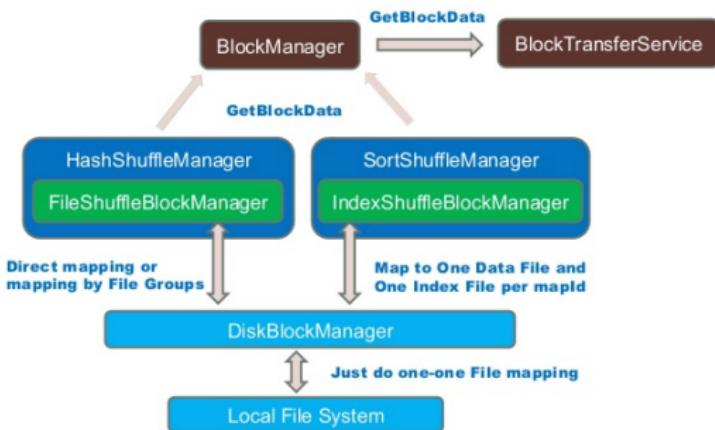


• Data Aggregation

- ▶ Defined on `ShuffleMapTask`
- ▶ Two methods available:
 - ★ `AppendOnlyMap`: in-memory hash table combiner
 - ★ `ExternalAppendOnlyMap`: memory + disk hash table combiner

• Batching disk writes to increase throughput

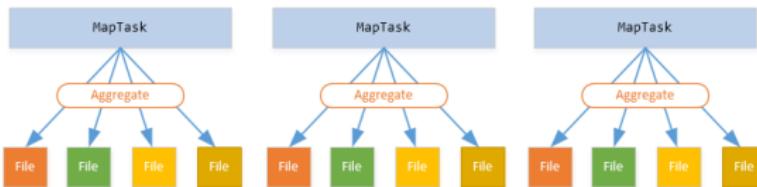
The Spark Shuffle Mechanism: Implementation Details



● Pluggable component

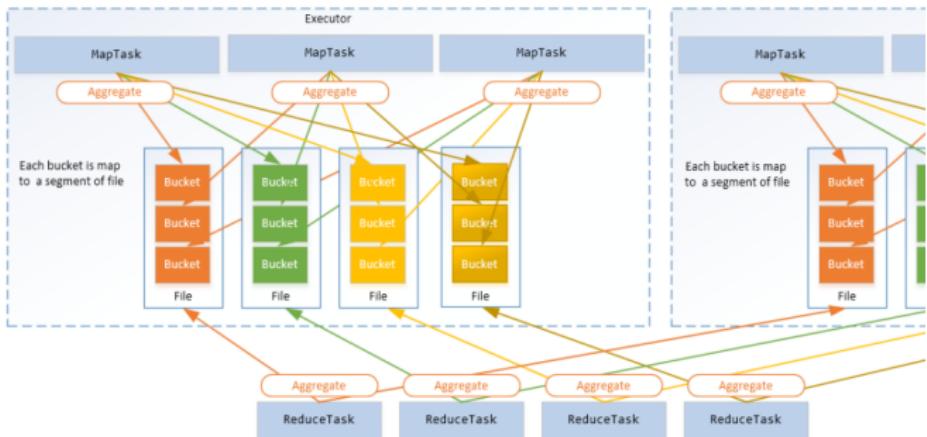
- ▶ **Shuffle Manager**: components registered to `SparkEnv`, configured through `SparkConf`
- ▶ **Shuffle Writer**: tracks “intermediate data” for the `MapOutputTracker`
- ▶ **Shuffle Reader**: pull-based mechanism used by the `ShuffleRDD`
- ▶ **Shuffle Block Manager**: mapping between logical partitioning and the physical layout of data

The Hash Shuffle Mechanism



- **Map Tasks write output to multiple files**
 - ▶ Assume: m map tasks and r reduce tasks
 - ▶ Then: $m \times r$ shuffle files as well as in-memory buffers (for batching writes)
- **Be careful on storage space requirements!**
 - ▶ Buffer size must not be too big with many tasks
 - ▶ Buffer size must not be too small, for otherwise throughput decreases

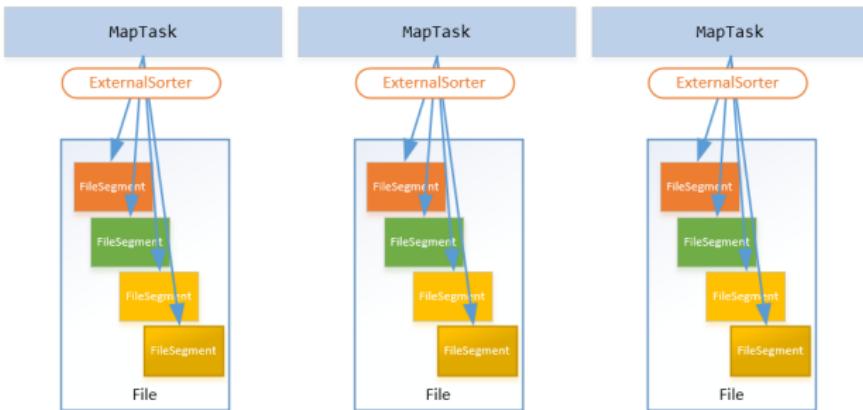
The Consolidate Hash Shuffle Mechanism



- Addresses buffer size problems

- ▶ Executor view vs. Task view
- ▶ Buckets are consolidated in a single file
- ▶ Hence: $F = C \times r$ files and buffers, where C is the number of Task threads within an Executor

The Sort-based Shuffle Mechanism



- **Implements the Hadoop Shuffle mechanism**
 - ▶ Single shuffle file, plus an index file to find “buckets”
 - ▶ Very beneficial for write throughput, as more disk writes can be batched
- **Sorting mechanism**
 - ▶ Pluggable external sorter
 - ▶ Degenerates to Hash Shuffle if no sorting is required

Data Transfer: Implementation Details

- **BlockTransfer Service**

- ▶ General interface for ShuffleFetcher
- ▶ Uses BlockDataManager to get local data

- **Shuffle Client**

- ▶ Manages and wraps the “client-side”, setting up the TransportContext **and** TransportClient

- **Transport Context:** manages the transport layer

- **Transport Server:** streaming server

- **Transport Client:** fetches consecutive chunks

Data Transfer: an Illustration

